

Operating Systems (SFWRENG 3SH3), Winter 2022

Prof. Neerja Mhaskar

Extra material for Assignment 1

/proc File System

For Kernel Versions <= 6.2.x (as per the textbook edition) please refer to the oldKernel folder, otherwise for newer Linux distributions please refer to the newKernel folder provided. To find the kernel version please run the following command on the terminal `uname -r`.

Note: hello_newKernel.c and hello_oldKernel.c is simply referred to as hello.c from here on to avoid duplication of material.

The /proc file system is a “pseudo” file system that exists only in kernel memory and is used primarily for querying various kernel and per-process statistics. This exercise involves designing kernel modules that create additional entries in the /proc file system involving both kernel statistics and information related to specific processes.

We begin by describing how to create a new entry in the /proc file system. The program example named hello.c (included with this PDF) creates a /proc entry named /proc/hello. If a user enters the command

```
cat /proc/hello
```

the Hello World message is returned.

Older Kernels (using struct file_operations):

In older kernels, we create a new /proc/hello entry using the proc_create() function in the module entry point proc_init(). The proc_create() function is passed a reference to a struct file_operations, which defines the operations that can be performed on the /proc/hello file. This structure initializes members such as .owner and .read. The .owner member is set to THIS_MODULE, which tracks the module that owns the /proc entry, while the .read member points to the proc_read() function. This function will be invoked whenever the /proc/hello file is read, ensuring that the module provides the correct behavior when userspace accesses the /proc entry.

Newer Kernels (using struct proc_ops):

In newer kernels, the `proc_create()` function uses a `struct proc_ops` instead of `struct file_operations` for handling /proc entries. The `proc_ops` structure replaces the older `file_operations` for /proc file system operations, providing a more focused set of operations specific to /proc file handling. The `.proc_read` member in `proc_ops` replaces `.read`, and it is similarly assigned to the `proc_read()` function. This change simplifies handling /proc files while keeping other file operations separate. As in older kernels, `proc_create()` still associates the module with the `/proc/hello` entry and ensures the `proc_read()` function is called when userspace reads from the file.

Examining this `proc_read()` function, we see that the string "Hello World\n" is written to the variable buffer where buffer exists in kernel memory. Since `/proc/hello` can be accessed from user space, we must copy the contents of buffer to user space using the kernel function `copy_to_user()`. This function copies the contents of kernel memory buffer to the variable `usr_buf`, which exists in user space.

Each time the `/proc/hello` file is read, the `proc_read()` function is called repeatedly until it returns 0, so there must be logic to ensure that this function returns 0 once it has collected the data (in this case, the string "Hello World\n") that is to go into the corresponding `/proc/hello` file.

Finally, notice that the `/proc/hello` file is removed in the module exit point `proc_exit()` using the function `remove_proc_entry()`.

Please use the following commands to try out the above hello kernel module.

1. `make`
2. `sudo insmod hello.ko` (`hello_oldKernel.ko` or `hello_newKernel.ko`)
3. `cat /proc/hello`
4. `sudo rmmod hello` (`hello_oldKernel` or `hello_newKernel`)